

Introduction to dplyr and magrittr

Denver R Users Group

www.meetup.com/DenverRUG

Peter DeWitt

peter.dewitt@ucdenver.edu

1 July 2014

Goals:

- ▶ Showcase dplyr, compare the ease of use compared to base R.
- ▶ Introduce the data manipulation grammar and philosophy behind dplyr
- ▶ Illustrate the usefulness of the forward-piping operator which is part of dplyr and extended further in magrittr.

dplyr

Data Import

dplyr verbs

select

arrange

filter

mutate

summarize

group_by

Chaining Work together

Joins

Memory Usage

Window Functions

Output

dplyr: a grammar of data manipulation

- ▶ Authored by Hadley Wickham and Romain Francois
- ▶ Current CRAN version 0.2

dplyr: a grammar of data manipulation

- ▶ Authored by Hadley Wickham and Romain Francois
- ▶ Current CRAN version 0.2
- ▶ Paraphrasing from a post on the RStudio blog <http://blog.rstudio.org/2014/01/17/introducing-dplyr>
 - ▶ dplyr is the next iteration of plyr
 - ▶ focuses only on `data.frames`
 - ▶ faster, thanks in part to Francois work in Rcpp, some use of multiple processors.
 - ▶ improved API.
 - ▶ interface with remote database (PostgreSQL, MySQL, SQLite, and Google bigquery) tables using the same verbs for interacting with `data.frames`. (Extendible to other backends)
 - ▶ Common operations:
 - ▶ `group_by`, `summarize`, `mutate`, `filter`, `select`, and `arrange`.

Data Import

dplyr does not have special tools for reading in data, but, if you need to rbind sets together...

```
# FAAs wildlife strikes on aircraft since 1990. The data  
# can be downloaded, in a Microsoft Access DB, from  
# http://www.faa.gov/airports/airport_safety/wildlife/database/  
# Tables in the DB were exported to csv files.  
# A data dictionary, in an Excel file, was also  
# included in the download from faa.gov
```

```
# column classes are set (in R code not shown) to ensure  
# that each column of the imported data is of the same class
```

```
wls.90.99 <-
```

```
  read.csv("../data/STRIKE_REPORTS (1990-1999).csv",  
            colClasses = clclass)
```

```
wls.00.09 <-
```

```
  read.csv("../data/STRIKE_REPORTS (2000-2009).csv",  
            colClasses = clclass)
```

```
wls.10.14 <-
```

```
  read.csv("../data/STRIKE_REPORTS (2010-Current).csv",  
            colClasses = clclass)
```

Data Import

```
# Base does not require the columns to be of the same class,  
# only the same name  
# dplyr requires that the columns are of the same class.  
dim(wls.90.99)  
  
## [1] 30150      94  
  
nrow(wls.90.99) + nrow(wls.00.09) + nrow(wls.10.14)  
  
## [1] 142911  
  
bnchmrk <-  
  benchmark(base = rbind(wls.90.99, wls.00.09, wls.10.14),  
            dplyr = rbind_list(wls.90.99, wls.00.09, wls.10.14),  
            replications = 100)  
bnchmrk[, c("test", "replications", "elapsed", "relative")]  
  
##      test replications elapsed relative  
## 1  base           100    92.98    3.968  
## 2 dplyr           100    23.43    1.000
```

Data Import

```
wls_df <- rbind(wls.90.99, wls.00.09, wls.10.14)
class(wls_df)

## [1] "data.frame"

wls <- rbind_list(wls.90.99, wls.00.09, wls.10.14)
class(wls)

## [1] "data.frame"

# A data frame tbl wraps a local data frame. The main
# advantage to using a tbl_df over a regular data frame is
# the printing: tbl objects only print a few rows and all
# the columns that fit on one screen, providing describing
# the rest of it as text. [source: R help doc]
wls_tbl_df <- tbl_df(wls)
class(wls_tbl_df)

## [1] "tbl_df"      "tbl"        "data.frame"
```

Data Printing

```
# print(wls_df)  # takes a long time, not helpful
# head(wls_df)   # too many columns to be useful
print(wls_tbl_df, n = 2)
```

```
## Source: local data frame [142,911 x 94]
```

```
##
```

```
##      INDEX_NR OPID      OPERATOR      ATYPE AMA AMO EMA EMO
## 1      100000  AAL  AMERICAN AIRLINES      B-727 148  10  34  10
## 2      100001  UAL   UNITED AIRLINES B-737-300 148  24  10  01
## ..      ...   ...                ...      ...  ...  ...  ...
```

```
## Variables not shown: AC_CLASS (chr), AC_MASS (int), NUM_ENGS
## (chr), TYPE_ENG (chr), ENG_1_POS (chr), ENG_2_POS (int),
## ENG_3_POS (chr), ENG_4_POS (int), REG (chr), FLT (chr),
## REMAINS_COLLECTED (lgl), REMAINS_SENT (lgl), INCIDENT_DATE
## (chr), INCIDENT_MONTH (int), INCIDENT_YEAR (int),
## TIME_OF_DAY (chr), TIME (int), AIRPORT_ID (chr), AIRPORT
## (chr), STATE (chr), FAAREGION (chr), ENROUTE (chr), RUNWAY
## (chr), LOCATION (chr), HEIGHT (int), SPEED (int), DISTANCE
## (dbl), PHASE_OF_FLT (chr), DAMAGE (chr), STR_RAD (lgl),
## DAM_RAD (lgl), STR_WINDSHLD (lgl), DAM_WINDSHLD (lgl),
## STR_NOSE (lgl), DAM_NOSE (lgl), STR_ENG1 (lgl), DAM_ENG1
## (lgl), STR_ENG2 (lgl), DAM_ENG2 (lgl), STR_ENG3 (lgl),
```


The verbs

- ▶ “Variable and function names should be lowercase. Use an underscore (_) to separate words within a name. Generally, variable names should be nouns and function names should be verbs. Strive for names that are concise and meaningful (this is not easy!).” - Hadley Wickham,
<http://adv-r.had.co.nz/Style.html>
- ▶ Verbs in dplyr
 - ▶ select,
 - ▶ arrange,
 - ▶ filter,
 - ▶ mutate,
 - ▶ summarize.

select

```
# Select columns of a data.frame, tbl_df.
wls_yr <- select(wls_tbl_df, INCIDENT_YEAR, AIRPORT,
                  ENG_1_POS, ENG_2_POS, DAM_ENG1, DAM_ENG2,
                  HEIGHT, DISTANCE, SPEED)
print(wls_yr, n = 5)

## Source: local data frame [142,911 x 9]
##
##   INCIDENT_YEAR      AIRPORT ENG_1_POS
## 1      1992 DALLAS/FORT WORTH INTL ARPT      5
## 2      1996      SACRAMENTO INTL          1
## 3      1996      DENVER INTL AIRPORT      1
## 4      1996      EPPLEY AIRFIELD         1
## 5      1996 WASHINGTON DULLES INTL ARPT      1
## ..          ...          ...          ...
## Variables not shown: ENG_2_POS (int), DAM_ENG1 (lgl),
##   DAM_ENG2 (lgl), HEIGHT (int), DISTANCE (dbl), SPEED (int)
```

select

```
# relative speed
bnch <-
  benchmark(base = wls_tbl_df[, c("INCIDENT_YEAR", "AIRPORT",
                                   "ENG_1_POS", "ENG_2_POS",
                                   "DAM_ENG1", "DAM_ENG2",
                                   "HEIGHT", "DISTANCE", "SPEED")],
            dplyr = select(wls_tbl_df,
                           INCIDENT_YEAR, AIRPORT,
                           ENG_1_POS, ENG_2_POS,
                           DAM_ENG1, DAM_ENG2,
                           HEIGHT, DISTANCE, SPEED),
            replications = 100)
select(bnch, test, replications, elapsed, relative)

##      test replications elapsed relative
## 1  base             100   0.007    1.000
## 2 dplyr             100   0.032    4.571
```

Selection of columns might be slower, but, there are some tools to help speed up the coding, and maintenance.

select

```
# num_range("x", 1:5, width = 2): selects all variables  
# (numerically) from x01 to x05.
```

```
select(wls_tbl_df, num_range("DAM_ENG", 1:4))
```

```
## Source: local data frame [142,911 x 4]
```

```
##
```

```
##      DAM_ENG1 DAM_ENG2 DAM_ENG3 DAM_ENG4
```

```
## 1      FALSE      FALSE      FALSE      FALSE
```

```
## 2      FALSE      FALSE      FALSE      FALSE
```

```
## 3      FALSE      FALSE      FALSE      FALSE
```

```
## 4      FALSE      FALSE      FALSE      FALSE
```

```
## 5      FALSE      FALSE      FALSE      FALSE
```

```
## 6      FALSE      FALSE      FALSE      FALSE
```

```
## 7      FALSE      FALSE      FALSE      FALSE
```

```
## 8      FALSE      FALSE      FALSE      FALSE
```

```
## 9      FALSE      FALSE      FALSE      FALSE
```

```
## 10     FALSE      FALSE      FALSE      FALSE
```

```
## ..      ...      ...      ...      ...
```

select

```
# starts_with(x, ignore.case = FALSE): names starts with x
select(wls_tbl_df, starts_with("DAM"))
```

```
## Source: local data frame [142,911 x 15]
```

```
##
```

```
##      DAMAGE DAM_RAD DAM_WINDSHLD DAM_NOSE DAM_ENG1 DAM_ENG2
```

```
## 1          N   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 2          FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 3          FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 4          N   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 5          N   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 6          M   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 7          N   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 8         M?   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 9          N   FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## 10         FALSE          FALSE    FALSE    FALSE    FALSE
```

```
## ..      ...      ...      ...      ...      ...      ...
```

```
## Variables not shown: DAM_ENG3 (lgl), DAM_ENG4 (lgl), DAM_PROP
```

```
##      (lgl), DAM_WING_ROT (lgl), DAM_FUSE (lgl), DAM_LG (lgl),
```

```
##      DAM_TAIL (lgl), DAM_LGHTS (lgl), DAM_OTHER (lgl)
```

select

```
# ends_with(x, ignore.case = FALSE): names ends in x  
select(wls_tbl_df, ends_with("4"))
```

```
## Source: local data frame [142,911 x 2]
```

```
##
```

```
##      STR_ENG4  DAM_ENG4
```

```
## 1      FALSE    FALSE
```

```
## 2      FALSE    FALSE
```

```
## 3      FALSE    FALSE
```

```
## 4      FALSE    FALSE
```

```
## 5      FALSE    FALSE
```

```
## 6      FALSE    FALSE
```

```
## 7      FALSE    FALSE
```

```
## 8      FALSE    FALSE
```

```
## 9      FALSE    FALSE
```

```
## 10     FALSE    FALSE
```

```
## ..      ...      ...
```

select

```
# matches(x, ignore.case = FALSE): selects all variables
# whose name matches the regular expression x
select(wls_tbl_df, matches("ENG|DAM"))
```

```
## Source: local data frame [142,911 x 26]
```

```
##
```

| ## | NUM_ENGS | TYPE_ENG | ENG_1_POS | ENG_2_POS | ENG_3_POS | ENG_4_POS |
|-------|----------|----------|-----------|-----------|-----------|-----------|
| ## 1 | 3 | D | 5 | 6 | 5 | NA |
| ## 2 | 2 | D | 1 | 1 | | NA |
| ## 3 | 2 | D | 1 | 1 | | NA |
| ## 4 | 2 | D | 1 | 1 | | NA |
| ## 5 | 2 | D | 1 | 1 | | NA |
| ## 6 | 2 | D | 1 | 1 | | NA |
| ## 7 | 3 | D | 5 | 6 | 5 | NA |
| ## 8 | 2 | C | 4 | 4 | | NA |
| ## 9 | 2 | D | 1 | 1 | | NA |
| ## 10 | 2 | D | 5 | 5 | | NA |
| ## .. | ... | ... | ... | ... | ... | ... |

```
## Variables not shown: DAMAGE (chr), DAM_RAD (lgl),
## DAM_WINDSHLD (lgl), DAM_NOSE (lgl), STR_ENG1 (lgl),
## DAM_ENG1 (lgl), STR_ENG2 (lgl), DAM_ENG2 (lgl), STR_ENG3
## (lgl), DAM_ENG3 (lgl), STR_ENG4 (lgl), DAM_ENG4 (lgl),
```

select

```
# contains(x, ignore.case = FALSE): selects all
```

```
# variables whose name contains x
```

```
select(wls_tbl_df, contains("ENG"))
```

```
## Source: local data frame [142,911 x 14]
```

```
##
```

```
##      NUM_ENGS TYPE_ENG ENG_1_POS ENG_2_POS ENG_3_POS ENG_4_POS
## 1           3         D          5          6          5         NA
## 2           2         D          1          1          1         NA
## 3           2         D          1          1          1         NA
## 4           2         D          1          1          1         NA
## 5           2         D          1          1          1         NA
## 6           2         D          1          1          1         NA
## 7           3         D          5          6          5         NA
## 8           2         C          4          4          1         NA
## 9           2         D          1          1          1         NA
## 10          2         D          5          5          1         NA
## ..          ...         ...         ...         ...         ...         ...
```

```
## Variables not shown: STR_ENG1 (lg1), DAM_ENG1 (lg1), STR_ENG2
```

```
## (lg1), DAM_ENG2 (lg1), STR_ENG3 (lg1), DAM_ENG3 (lg1),
```

```
## STR_ENG4 (lg1), DAM_ENG4 (lg1)
```


select

What about dropping variables?

```
print(wls_yr, n = 2)

## Source: local data frame [142,911 x 9]
##
##      INCIDENT_YEAR      AIRPORT ENG_1_POS
## 1      1992 DALLAS/FORT WORTH INTL ARPT      5
## 2      1996      SACRAMENTO INTL          1
## ..      ...                ...      ...
## Variables not shown: ENG_2_POS (int), DAM_ENG1 (lgl),
##      DAM_ENG2 (lgl), HEIGHT (int), DISTANCE (dbl), SPEED (int)
```

```
print(select(wls_yr, -AIRPORT, -starts_with("ENG")), n = 3)
```

```
## Source: local data frame [142,911 x 6]
##
##      INCIDENT_YEAR DAM_ENG1 DAM_ENG2 HEIGHT DISTANCE SPEED
## 1      1992      FALSE      FALSE    300      NA    142
## 2      1996      FALSE      FALSE      0      0     NA
## 3      1996      FALSE      FALSE      0      0     NA
## ..      ...                ...      ...      ...      ...
```

arrange

arrange: reorder the rows. Multiple inputs are ordered from left-to-right.

```
dat <- data.frame(var1 = c(3, 8, 2, 1),  
                  var2 = c("E", "A", "A", "B"))
```

```
dat
```

```
##   var1 var2  
## 1    3    E  
## 2    8    A  
## 3    2    A  
## 4    1    B
```

```
# this would be very helpful for collecting data by a  
# subject id, visit number, ...
```

arrange

```
arrange(dat, var2)
```

```
##   var1 var2
## 1    8    A
## 2    2    A
## 3    1    B
## 4    3    E
```

```
arrange(dat, var2, var1)
```

```
##   var1 var2
## 1    2    A
## 2    8    A
## 3    1    B
## 4    3    E
```

*# this would be very helpful for collecting data by a
subject id, visit number, ...*

filter

filter: return only a subset of the rows. If multiple conditions are supplied they are combined with &.

```
dim(wls_yr)

## [1] 142911      9

filter(wls_yr, INCIDENT_YEAR > 2000, INCIDENT_YEAR <= 2005)

## Source: local data frame [31,947 x 9]
##
##   INCIDENT_YEAR                                AIRPORT ENG_1_POS
## 1             2001                JOHN F KENNEDY INTL
## 2             2001            SAN FRANCISCO INTL ARPT      1
## 3             2001                ORLANDO INTL            1
## 4             2001                MOLOKAI ARPT            4
## 5             2001            LAMBERT-ST LOUIS INTL        5
## 6             2001                KANSAS CITY INTL        1
## 7             2001                UNKNOWN                 1
## 8             2001            AKRON-CANTON REGIONAL        7
## 9             2001 DESTIN-FORT WALTON BEACH ARPT          5
## 10            2001                JOHN F KENNEDY INTL      1
##
```

filter

```
bnch <-  
  benchmark(base = subset(wls_yr, INCIDENT_YEAR > 2000 & INCIDENT_YEAR  
    dplyr = filter(wls_yr, INCIDENT_YEAR > 2000, INCIDENT_YEAR  
      replications = 100)  
select(bnch, test, replications, elapsed, relative)  
  
##      test replications elapsed relative  
## 1  base             100   8.038     5.402  
## 2 dplyr             100   1.488     1.000
```

mutate

mutate: add new columns. Multiple inputs create multiple columns.

```
eng.lbls <- c("mounted below the wing", "mounted above the wing",
              "part of the wing root", "nacelle-mounted on the wing",
              "mounted on the aft fuselage")

str(mutate(wls_yr,
           SPEED_MPH = SPEED * 1.15078, # SPEED was in knots
           ENG_1_POS = factor(ENG_1_POS, 19:23, eng.lbls),
           ENG_2_POS = factor(ENG_2_POS, 19:23, eng.lbls)))

## Classes 'tbl_df', 'tbl' and 'data.frame': 142911 obs. of  10 variables:
## $ INCIDENT_YEAR: int  1992 1996 1996 1996 1996 1996 1991 1993 1995
## $ AIRPORT      : chr  "DALLAS/FORT WORTH INTL ARPT" "SACRAMENTO INTL
## $ ENG_1_POS    : Factor w/ 5 levels "mounted below the wing",... NA
## $ ENG_2_POS    : Factor w/ 5 levels "mounted below the wing",... NA
## $ DAM_ENG1     : logi  FALSE FALSE FALSE FALSE FALSE FALSE ...
## $ DAM_ENG2     : logi  FALSE FALSE FALSE FALSE FALSE FALSE ...
## $ HEIGHT       : int   300 0 0 0 1000 5000 0 1500 0 100 ...
## $ DISTANCE     : num   NA 0 0 0 NA NA 0 NA 0 NA ...
## $ SPEED        : int   142 NA NA NA NA NA 100 220 NA 135 ...
## $ SPEED_MPH    : num   163 NA NA NA NA ...
```

mutate

```
bnch <-  
benchmark(base = within(wls_yr, {  
  SPEED_MPH = SPEED * 1.15078  
  ENG_1_POS = factor(ENG_1_POS, 19:23, eng.lbls)  
  ENG_2_POS = factor(ENG_2_POS, 19:23, eng.lbls)}  
  dplyr = mutate(wls_yr,  
    SPEED_MPH = SPEED * 1.15078,  
    ENG_1_POS = factor(ENG_1_POS, 19:23, eng.lbls)  
    ENG_2_POS = factor(ENG_2_POS, 19:23, eng.lbls)  
    replications = 100)  
select(bnch, test, replications, elapsed, relative)
```

| ## | test | replications | elapsed | relative |
|------|-------|--------------|---------|----------|
| ## 1 | base | 100 | 5.824 | 1.067 |
| ## 2 | dplyr | 100 | 5.458 | 1.000 |

summarize

summarise: reduce each group to a single row. Multiple inputs create multiple output summaries. (Two spellings: summarize and summarise.)

```
summarise(wls_yr,  
  "Mean speed" = mean(SPEED, na.rm = TRUE),  
  "SD speed"   = sd(SPEED, na.rm = TRUE),  
  n            = sum(!is.na(SPEED)))
```

```
## Source: local data frame [1 x 3]
```

```
##
```

```
##   Mean speed SD speed      n
```

```
## 1      141.3   46.09 58938
```


group_by

```
summarise(group_by(wls_yr, ENG_1_POS),  
           "Mean speed" = mean(SPEED, na.rm = TRUE),  
           "SD speed"   = sd(SPEED, na.rm = TRUE),  
           n             = sum(!is.na(SPEED)))
```

```
## Source: local data frame [11 x 4]
```

```
##
```

| ## | ENG_1_POS | Mean speed | SD speed | n |
|-------|-----------|------------|----------|-------|
| ## 1 | | 113.20 | 40.40 | 1303 |
| ## 2 | 1 | 154.81 | 43.14 | 27634 |
| ## 3 | 2 | 62.00 | 43.39 | 4 |
| ## 4 | 3 | 108.32 | 33.48 | 31 |
| ## 5 | 4 | 123.90 | 41.74 | 7953 |
| ## 6 | 5 | 143.80 | 42.31 | 17701 |
| ## 7 | 6 | 99.06 | 36.67 | 482 |
| ## 8 | 7 | 83.91 | 29.34 | 3829 |
| ## 9 | A | 90.00 | NA | 1 |
| ## 10 | C | NaN | NA | 0 |
| ## 11 | T | NaN | NA | 0 |

group_by

```
bnch <-  
  benchmark(base = aggregate(SPEED ~ ENG_1_POS, wls_yr,  
                             function(x) c(mean = mean(x, na.rm = TRUE),  
                                             sd   = sd(x, na.rm = T),  
                                             n    = sum(!is.na(x)))),  
            dplyr = summarise(group_by(wls_yr, ENG_1_POS),  
                              "Mean speed" = mean(SPEED, na.rm = TRUE),  
                              "SD speed"   = sd(SPEED, na.rm = TRUE),  
                              n            = sum(!is.na(SPEED))),  
            replications = 100)  
select(bnch, test, replications, elapsed, relative)
```

```
##      test replications elapsed relative  
## 1  base           100  79.457    49.72  
## 2 dplyr           100   1.598     1.00
```

Say we need to filter, group_by, and summarise data

*# What is the mean distance from the airport, in kilometers,
where the strike took place, by damage to engine, on twin
engine aircraft, between 2002 and 2010, inclusive?*

```
summarize(group_by(mutate(filter(wls, INCIDENT_YEAR >= 2002, INCIDENT_YEAR <= 2010, NUM_ENGS == 2), DISTANCE_KM = DISTANCE * 1.60934), DAM_ENG1, DAM_ENG2), `mean distance in KM` = mean(DISTANCE_KM, na.rm = TRUE))
```

```
## Source: local data frame [4 x 3]
```

```
## Groups: DAM_ENG1
```

```
##
```

| ## | DAM_ENG1 | DAM_ENG2 | mean distance in KM |
|------|----------|----------|---------------------|
| ## 1 | FALSE | FALSE | 1.3715 |
| ## 2 | FALSE | TRUE | 1.3228 |
| ## 3 | TRUE | FALSE | 0.8347 |
| ## 4 | TRUE | TRUE | 0.6584 |

*# Without a comment to explain, how long would it take to
explain the above code? You need to read from the inside
out. THERE IS A BETTER WAY*

Chain together multiple operations.

```
wls %>%
  filter(INCIDENT_YEAR >= 2002,
         INCIDENT_YEAR <= 2010,
         NUM_ENGS == 2) %>%
  mutate(DISTANCE_KM = DISTANCE * 1.60934) %>%
  group_by(DAM_ENG1, DAM_ENG2) %>%
  summarise("mean distance in KM" = mean(DISTANCE_KM, na.rm = TRUE))

## Source: local data frame [4 x 3]
## Groups: DAM_ENG1
##
##   DAM_ENG1 DAM_ENG2 mean distance in KM
## 1   FALSE   FALSE      1.3715
## 2   FALSE   TRUE      1.3228
## 3    TRUE   FALSE      0.8347
## 4    TRUE   TRUE      0.6584
```

More detailed examples of the forward-piping operator follow.

joining data sets

- ▶ dplyr version 0.2 has the following joins:
 - ▶ `inner_join`,
 - ▶ `left_join`,
 - ▶ `semi_join`, and
 - ▶ `anti_join`.
- ▶ Stated milestone for version 0.3 includes
 - ▶ `outer_join`,
 - ▶ `right_join`, and
 - ▶ `cross_join`.

Improved memory usage

```
this_df2 <- this_df <- data.frame(var1 = 1:5, var2 = rnorm(5))
changes(this_df, this_df2)

## <identical>

this_df$var1 <- rexp(5, rate = 2)
changes(this_df, this_df2)

## Changed variables:
##           old           new
## var1      0x157ae870 0x62dca90
##
## Changed attributes:
##           old           new
## row.names 0x1640b390 0x157bd6b0
```

Improved memory usage

```
this_df2 <- this_df <- data.frame(var1 = 1:5, var2 = rnorm(5)) %>% tbl_
changes(this_df, this_df2)
```

```
## <identical>
```

```
this_df <- this_df %>% mutate(var1 = rexp(5))
changes(this_df, this_df2)
```

```
## Changed variables:
```

```
##           old           new
## var1      0x1936de38 0xd580f48
##
```

```
## Changed attributes:
```

```
##           old           new
## names      0x87e4618 0xc1e0958
## class      0x18ed2ad0 0xd580980
## row.names  0x4958c40 0x4958fa0
```

dplyr memory usage

From the vignette("memory", "dplyr")

- ▶ `tbl_df()` and `group_by()` don't copy columns
- ▶ `select()` never copies columns, even when you rename them
- ▶ `mutate()` never copies columns, except when you modify an existing column
- ▶ `arrange()` must copy because you're changing the order of every column. This is an expensive operation for big data, but you can generally avoid it using the `order` argument to window functions
- ▶ `summarise()` creates new data, but it's usually at least an order of magnitude smaller than the original data.

Window Functions

```
# see vignette("window-functions", package = "dplyr")
```

Other Data Sources

- ▶ `dplyr` works for
 - ▶ `data.frames`,
 - ▶ `data.tables`, databases, and multidimensional arrays.
 - ▶ Same verbs used for all data sources.
 - ▶ See `vignette("databases", package = "dplyr")` for more details.

data.table vs dplyr

From the
dplyr introduction vignette:

- ▶ For multiple operations, data.table can be faster because you usually use it with multiple verbs at the same time. For example, with data table you can do a mutate and a select in a single step, and it's smart enough to know that there's no point in computing the new variable for the rows you're about to throw away.
- ▶ The advantages of using dplyr with data tables are:
 - ▶ For common data manipulation tasks, it insulates you from reference semantics of data.tables, and protects you from accidentally modifying your data.
 - ▶ Instead of one complex method built on the subscripting operator (`[]`), it provides many simple methods.

magrittr: a forward-pipe operator for R

ceci n'est pas un pipe (this is not a pipe)

- ▶ dplyr functionality is made more powerful via the `%>%`, or equivalently, `\%.%$`, operator.
- ▶ Additional functionality provided by the `magrittr` package authored by Stefan Bache and Hadley Wickham.
- ▶ These operators are similar to
 - ▶ F#'s `|>`, or
 - ▶ Linux's `|`.
- ▶ Use of these operators will drastically change your R syntax.
- ▶ Helpful to writing complex, nested, operations.
- ▶ “Read from left to right instead of inside out.”

magrittr: a forward-pipe operator for R

Examples

```
data(diamonds, package = "ggplot2")

# find the mean price of the diamonds
# Standard R syntax
mean(diamonds$price)

## [1] 3933

# with the pipe
diamonds %>%
  extract("price") %>%
  unlist() %>%
  mean()

## [1] 3933
```

What's the point?

Reproducibility

The data, code, sides, etc. all at
github.com/dewittpe/dplyr-demo

```
print(sessionInfo(), locale = FALSE)

## R version 3.1.0 (2014-04-10)
## Platform: x86_64-pc-linux-gnu (64-bit)
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets
## [6] methods    base
##
## other attached packages:
## [1] rbenchmark_1.0.0 dplyr_0.2      magrittr_1.0.1
## [4] knitr_1.6         vimcom_0.9-93  setwidth_1.0-3
## [7] colorout_1.0-3
##
## loaded via a namespace (and not attached):
## [1] assertthat_0.1  codetools_0.2-8 digest_0.6.4
## [4] evaluate_0.5.5  formatR_0.10    highr_0.3
## [7] parallel_3.1.0  qwraps_0.2.2    Rcpp_0.11.1
## [10] stringr_0.6.2   tcltk_3.1.0     tools_3.1.0
```

DRUG

- ▶ Future MeetUp Topics:
 - ▶ (Possible) iPython / R speaker for later in July
 - ▶ We need others speakers!
- ▶ MeetUp locations/times